



AI and the Humanities: A Framework for Language and Literary Scholarship

Prepared by the [AI and Research Working Group](#)

June 2026

AI and Research Working Group

Maria Antoniak (University of Colorado, Boulder)
Nina Beguš (University of California, Berkeley)
Kathi Inman Berens (Portland State University)
André Brock (Georgia Institute of Technology)
Xuenan Cao (Chinese University of Hong Kong)
Kirsti Cole (North Carolina State University)
Leonardo Flores (Appalachian State University); cochair
Md Mahmudul Haque (North Dakota State University)
Matthew Kirschenbaum (University of Virginia)
Lauren Klein (Emory University)
Kari Kraus (University of Maryland, College Park)
Liz Losh (William and Mary)
Meredith Martin (Princeton University)
David Mimno (Cornell University)
John Murray (University of Central Florida)
Kate Ozment (Case Western Reserve University)
Rita Raley (University of California, Santa Barbara)
Lisa Marie Rhody (Graduate Center, City University of New York); cochair
Jeffrey Tharsen (University of Chicago)
Lauren Tilton (University of Richmond)
Annette Vee (University of Pittsburgh)
Melanie Walsh (University of Washington, Seattle)

AI and the Humanities: A Framework for Language and Literary Scholarship

Introduction

Artificial intelligence (AI)* is transforming the disciplines of language and literary studies. Scholars read differently when they use AI to process texts at scales exceeding individual human capacity. The distinction between analysis and writing blurs when AI-generated prose is incorporated into iterative cycles of drafting and revision. When language models produce plausible, even sophisticated, content in areas where the writer lacks deep knowledge, the risk of uncritical acceptance increases and can compromise academic integrity. By amplifying existing tensions around collaboration and attribution, AI compels scholars to explicitly articulate—and critically reevaluate—the often-implicit norms around authorship, credit, and intellectual responsibility.

Recognizing that fundamental disciplinary commitments to interpretation, evidentiary accountability, and equitable scholarly labor are at stake, the Modern Language Association (MLA) convened the Task Force on AI in Research and Teaching to develop shared guidance for the profession. The task force then formed the AI and Research Working Group (AIRWG)—composed of scholars from diverse institutional contexts, regions, career paths, and fields of specialization—so that the guidance would reflect varied perspectives and modes of response.

This framework is designed to help the profession navigate new technical and social conditions. We encourage faculty members, administrators, editors, and publishers to use it as common ground for dialogue, consensus-building, and policy creation.

Framework for AI Use in Language and Literary Scholarship

This framework emerges from three interrelated areas—AI literacy, human-centered expertise, and responsibility and ethics—and is organized into thematic questions to guide discussion and critical reflection. Since no single policy can address every scenario, we demonstrate the framework’s utility through purposefully complex case studies (Appendices [A](#) and [B](#)). This heuristic approach acknowledges the

* In the context of this report, *AI* refers to generative AI tools such as *ChatGPT*, *Gemini*, *Claude*, and their many variants, including task-specific integrations embedded in research, writing, and coding environments. These are models and systems trained on vast amounts of data, often with chat interfaces, and used to generate text, images, code, and other media. Our focus is restricted to these generative technologies (primarily large language models), distinct from the machine learning, statistical, and database methods central to digital humanities and cognate practices.

messy reality of research, affirms the value of dialogue over dictates, and aspires to contribute to a larger conversation about AI use in the humanities.

AI Literacy

AI literacy is the critical foundation for the informed and ethical use of AI in humanities research and pedagogy. It moves beyond operational skill to encompass scholars' critical awareness of their positions in relation to AI technologies, as well as their embedded connections to disciplinary methods, data, and broader social contexts. This literacy requires probing the often-opaque architectures of AI systems to discern their inherent biases, costs, and methodological assumptions. Scholars practice AI literacy through due diligence on any tool under consideration, which may include consulting with relevant experts—digital humanists, computer scientists, librarians, or research software engineers—to address gaps in technical or contextual knowledge. Consistent with humanistic scholarship, this process begins with applying rigorous source criticism to the tool itself.

Functionality: What specific AI tool or model are you using, and what is its intended purpose? How does its design—its architecture, training data, and versioning—affect what it can (and cannot) do? For which tasks is it reliable, and what are its known limitations or failures? How will you verify or audit its outputs? Where is your data processed and stored? Is it handled locally, through an institutional sandbox, or through commercial cloud services, and what privacy protections apply? Has this tool been evaluated in peer-reviewed research or by your institution for work like yours?

Context: What social and ethical concerns are associated with this tool (e.g., data privacy, copyright, environmental impacts, labor exploitation, or embedded biases)? How do its (corporate) foundations and design philosophy shape how you interact with it? In what ways does its use support or challenge established methodological norms in your discipline? Where do you perceive tension between its practical benefits and broader ethical or scholarly considerations?

Precedents and Practices: What institutional, departmental, or professional policies govern your AI use? Do publishers or funding bodies have guidelines that apply? What documentation exists for this specific tool, and how transparent is it about its function and limitations? Are there established precedents in your field for using AI in this way—and when such uses exist, how have scholars disclosed them? Who can you consult for guidance (e.g., colleagues, librarians, research integrity officers, or technical specialists)? How will you ensure your approach reflects shared standards?

Human-Centered Expertise

Human-centered expertise draws on the interplay of AI literacy, disciplinary expertise, and lived experience. Though expertise has historically excluded knowledge gained through experience, rather than formal training, we redeploy the term *expertise* here to prioritize the situated understanding researchers develop through practice. As AI automates tasks once performed by humans, a human-centered model requires justifying the tool's use and safeguarding the scholar's role in feedback, iteration, and evaluation. This

section urges researchers to anchor their work in human agency and expertise, from initial inquiry to final communication.

Process: Why introduce AI at this stage of your research or writing? What specific challenge or need does it address? In what ways does it extend, replace, or reshape your interpretive agency, expertise, or creative labor? Is it a tool for exploration, or does it actively shape your argument, evidence, or rhetorical voice? How does its use reconfigure traditional phases of scholarly work? Crucially, do you retain the disciplinary fluency to assess, critique, and assume responsibility for what it produces?

Communication: How does using AI affect your collaboration with colleagues, coauthors, or subject-area experts? Does it support, substitute for, or complicate these vital scholarly relationships? In what ways does it shift your connection to readers, critics, the public, and the historical and intellectual record? When AI-generated text is paraphrased, adapted, or absorbed into your work, what happens to authorship, citation, and intellectual credit? How does this practice challenge (or potentially undermine) the communal, dialogic foundations of knowledge-making in your field?

Affect and Experience: What is the experience of using AI in your scholarly practice? Does it inspire creativity, induce anxiety, foster dependence, or create a sense of alienation from your own voice and labor? Whose cultural, historical, and linguistic perspectives shape the tool's outputs, and whose are absent or distorted? How might these biases subtly influence your interpretations or conclusions? In what ways does this technology center or marginalize particular forms of knowledge? Can human-in-the-loop practices help correct imbalances and amplify underrepresented voices? How does using AI reshape your understanding of scholarly authority, originality, and humanistic inquiry?

Responsibility and Ethics

Issues of ethics and integrity are the bedrock of research, yet AI complicates this foundation in unprecedented ways. Scholars must now grapple with challenges for which their training has not prepared them and at the same time attempt to reconcile technical capacities and humanistic values. This framework links responsibility to AI literacy and human-centered expertise, asking scholars to consider bias, privacy, and transparency as situated ethical problems, rather than technical roadblocks. Ultimately, these questions prompt a redefinition of responsible practice for institutions in the age of AI.

Bias: What assumptions do you bring to AI use, and how might they affect your work? In what ways can AI perpetuate bias, fabricate content, or blur lines of attribution? How will you remain accountable to ethical and scholarly standards in response?

Privacy: What data and prompts are you sharing with AI systems, and how might their use affect relationships with archival subjects, living communities, coauthors, or research participants? Could this compromise consent, confidentiality, or trust? How will you protect vulnerable information and uphold ethical responsibilities in collaborative or sensitive research contexts?

Transparency: When and why should AI use be documented in your work? Is transparency assumed, required, or contested by your field, publisher, or institution? How might AI obscure authorship or labor, and whose contributions risk being hidden or misrepresented? What disclosure practices (e.g., citation, footnotes, appendices) are necessary? If scholars use AI models to generate claims or insights, should they cite the human sources behind the model’s training data, just as they would a secondary source? How do legal frameworks like nondisclosure agreements or ethical requirements regarding conflicts of interest complicate these disclosures?

See [Appendix A: AI Use in Research](#) and [Appendix B: Assessing AI Use](#) for situations relevant to the questions presented above.

Conclusion

Living with Uncertainty: AI and the Future of Humanities Scholarship

The proliferation of AI technologies is occurring in the midst of other pressures facing humanities scholars, including limits on academic freedom, fiscal austerity, and attacks on organized labor and the institution of the university itself. By placing human expertise at the center of our work, we seek to resist the structural dehumanization of higher education—the fundamental objective of this document. However, AI tools are reshaping academic work faster than disciplinary norms and ethical frameworks can respond. Our goal, therefore, is to model an informed and critical stance that can continue to adapt as the technology and our understandings of its impacts develop.

In this context, we must clarify our commitments. What do we believe writing, reading, and thinking are for? What kinds of intellectual labor—drafting an argument, analyzing a text, engaging in creative expression—do we refuse to delegate to AI, even if it could do the work faster or more efficiently? And what forms of human attention, care, and judgment do we insist on preserving, not just as individual scholars but as a community? These questions are already being answered, not in mission and policy statements, but in everyday decisions about how we research, teach, and collaborate. Those decisions will define the humanities that emerge from this period and the publics they are able to serve.

What This Document Makes Possible

We hope this framework encourages multiple forms of productive conversation:

- **For individual scholars**, these questions provide a basis for reflective practice—a way to think through the goals of their research and the implications of AI use throughout a project’s lifecycle and to account for those choices in publications and presentations.
- **For departments and institutions**, this document can inform policy, professional development, and mentoring. Appendices [A](#) and [B](#) offer examples for collective deliberation, allowing communities to develop shared norms while respecting disciplinary differences.

- **For editors and reviewers**, the framework offers vocabulary and evaluative checkpoints for assessing manuscripts that were created with the assistance of AI or include AI-generated content. Rather than blanket acceptance or rejection, editors and reviewers can ask pertinent questions from the frameworks offered.
- **For graduate students and early career scholars**, who navigate particularly precarious labor conditions and face pressures to adopt efficiency tools, this document validates asking critical questions about institutional requirements and expectations.

Opportunities and Next Steps

The value of this framework will emerge through use, tests, and revision. We invite the MLA community to do the following:

- **Apply these questions** to your own research practices and share insights about where they prove helpful or fall short.
- **Adapt the scenarios** to your specific subfields, and develop new ones that illuminate further tensions.
- **Develop standardized disclosure practices.** A short process note can help concretize these norms without turning every manuscript into a methods appendix. For example:

AI Use Statement. This project used [tool/version] for [purposes] at [stages]. Outputs were evaluated for accuracy and bias, and all claims and citations derive from sources that were read and interpreted by the author(s), who accept(s) full responsibility for its content. No student or proprietary materials were entered without consent. Where language refinement occurred, the author(s) retained editorial control over meaning and style.

We encourage scholars to experiment with such statements in their publications and to share these efforts with colleagues and editors, who can refine and adopt them as emerging disciplinary conventions—much like the now-commonplace sharing of syllabus language regarding AI.

Finally, we recognize that individual scholars operate under varying degrees of institutional autonomy. Some may be required to use AI-integrated tools; others may have the opportunity to refuse; others might be incorporating these new methods yet fear retribution for disclosure of AI use. This document cannot resolve these structural differences, but it can provide language for articulating concerns, pushing for better institutional policies, and maintaining critical awareness even within constrained circumstances. In this new and fast-evolving environment, the continued free and open sharing of our ideas, methods, successes, and cautionary tales provides important guidance as we and our scholarly communities strive to assert the value of our work.

Appendix A: AI Use in Research

The following use cases are based on experiences that the authors of this paper have encountered in the current AI landscape, and we recognize that they may age. Nevertheless, we see them as valuable opportunities to reflect on issues that arise around the use of these technologies, and they serve neither as recommendations nor cautionary tales. Some people will be more comfortable with the uses described than others and may consider some aspects to be valuable and others a step too far. We invite you to think about these scenarios, discuss them with colleagues, and use them to guide reflections on your own potential use of AI. You may also wish to consider how your responses might differ if you are approaching them from the perspective of a graduate student, an early career scholar, an established faculty member, a department chair, or a dean.

Case #1 – Cursive Transcription

A researcher has cursive manuscripts that need transcription. The researcher has some funds, but not enough to pay an assistant to help with the transcriptions. A digital humanities specialist has helped them select an AI-powered tool that worked well on the manuscript and helped them set up a research workflow. The tool generates a readable transcript, and the researcher then checks the output, noting that the model can hallucinate and make mistakes that are different from those found in human-made transcripts. The AI model is not private; the data is sent to company servers and may be stored for future unknown uses (e.g., integration into training data).

Case #2 – Citing AI Output

A researcher has been working with an AI chatbot to help refine their understanding of a particular concept in their scholarly writing. They are careful and diligent, fact-checking the AI's output and using it as one component of their research process that also includes extensive engagement with published scholarship. At one point in a chat exchange, the AI produces a vivid and startling formulation of the concept that captures the essence of what they are trying to say. They want to use this language in a peer-reviewed article submitted to a top-tier humanities journal. Their instinct is that they should put the passage in quotation marks and cite it as they would any secondary source, but they don't know what the actual citation conventions are. They also worry this kind of transparent disclosure might prejudice the review of their article.

Case #3 – Generating an Outline and the Text of an Article

A scholar writes a detailed outline, including quotations and citations. They prompt an AI system to generate a draft, then iteratively refine it into a usable version. After copying it into a word processor, they revise and expand the draft until it is ready for peer-reviewed publication. The final version has been thoroughly revised to ensure accuracy, mitigate bias, offer a consistent style, and express the ideas the scholar wished to convey. They include a process note detailing their use of AI and their subsequent revisions.

Case #4 – Unwitting Use of AI in Language Translation

A scholar has conducted archival research and drawn up notes in the language used in the archive. In the process of drafting their research paper, they have used apps to translate the notes into another language. They then copied and pasted long paragraphs of their translated notes into the paper for submission, without knowing that the tools were AI-based. After the submission, the review process returned an AI detector's report accusing the scholar of including undisclosed AI-generated text. In retrospect, the scholar realized that the translation tool was the cause.

Case #5 – Presentation Transcription

A scholar delivers a conference presentation, which they audio-record. Afterward, they copy the voice-to-text transcript generated by the audio recording software into an AI tool and prompt it to organize the text into sentences and paragraphs. They paste this transcript into a word processor, develop it into a full essay, and submit it for publication. The essay includes a process note disclosing AI use in the transcription.

Case #6 – University Press Monograph That Incorporates AI-Assisted Research

A junior scholar up for tenure and promotion used AI tools in a data analytics framework for a large-scale corpus, resulting in a peer-reviewed, university press monograph that contributes novel work to their field. They know that vocal senior faculty members (in their own and neighboring departments) vehemently oppose AI on the basis that it is cheating, IP theft, ideologically suspect, and environmentally destructive. They also worry that reviewers will assume “the AI has done all the work” and are unlikely to be persuaded otherwise. They are unsure how to handle the situation, and how the chair and dean might help address it.

Case #7 – Using AI for a Book Review

A scholar has been invited to review a book for a journal that prohibits AI use in all circumstances. The scholar does not want to decline the request, but their scholarly workflow relies on AI tools (e.g., using an AI search interface to efficiently identify key or highly cited works in their field and cross-checking their chapter summaries against AI-generated summaries to spot omissions or biases). AI also assists them in drafting, editing, and organizing their writing. Given these circumstances, they are considering writing to the editors to seek an exception to the policy.

Appendix B: Assessing AI Use

The following scenarios describe situations in which you may encounter work that may or may not have used AI. They offer opportunities to reflect on issues that arise around the use of these technologies and are neither recommendations nor cautionary tales. Some people will be more comfortable with the uses described than others and may consider some aspects to be valuable and others a step too far. We invite you to think about these scenarios, discuss them with colleagues, and use them to guide reflections on

similar scenarios involving the use of AI. You may also wish to consider how your responses might differ if you are approaching them from the perspective of a graduate student, an early career scholar, an established faculty member, a department chair, or a dean.

Scenario #1 – Suspected AI-Generated Submission

You are a peer reviewer for a journal that publishes both creative and critical work and lacks a policy about AI use in submissions. One of the submitted works, a short story, reads as if AI generated, but it's hard to pinpoint exactly how. The prose is intriguing but seems a bit generic and uses rhetorical constructions that are popular with AI-generated text. Based on these characteristics, the editor asks whether you think this work may be AI generated and whether it is acceptable for publication.

Scenario #2 – Knowledge Gaps in an Essay Submission

You are editing a special issue for a journal in your field, and receive a submission by a single author new to your field. Reading the essay, you notice uncredited references to foundational ideas posed by major scholars. The author also coins terms and presents ideas established years ago by others, without attribution. Checking the author's prior publications (few and in a different field), you note a marked difference in writing style. The author has disclosed the use of a popular, free AI tool in the research, drafting, and editing of the essay, as required by the journal.

Scenario #3 – Digital Humanities Integrating AI

You agree to review a manuscript for an edited collection that is part of a digital humanities series by a reputable academic press. In it, there is a chapter in which a scholar discusses as part of their methodology that they used AI to generate code for topic modeling analysis of a large public domain textual corpus, also prepared with AI tools. They detail the use of a local, privacy-preserving AI system and the steps taken to protect against fabrication, hallucination, and bias. They further describe generating the code with a different, coding-optimized AI system. The full code is published on *GitHub* and linked in the manuscript.

Scenario #4 – AI Translation

You are editing a special issue on transnational literature for a peer-reviewed journal. One of the works submitted seems well researched and argued, but the prose style occasionally reads as if AI generated. Checking the scholar's publication history, you find they publish primarily in another language, and most cited sources cited are in that language, but note that the topic of the submitted work is well within their area of expertise. There is no process note or AI citation, though neither is required by the journal nor specified in the call for papers. You suspect AI was used to translate the work into English.

Scenario #5 – Considering AI Policies and Tools for a Journal

You are on the board of a journal discussing AI policies. The press asks if the journal should allow submitted content to be shared with AI models and companies. The editor in chief seeks input on the role of AI in the article writing process and peer review. An associate editor asks whether AI can be used to optimize or automate publication tasks (e.g., assigning peer reviewers). The press's technical team asks whether to integrate an AI detector into the content management system. Authors seek guidance on citing AI and disclosure statements.

Scenario #6 – AI in Tenure and Promotion Department-Level Guidelines

You are on a department-level tenure and promotion (T&P) committee evaluating candidates for advancement. Some members argue that AI use (e.g., in research central to domain knowledge) detracts from scholarship's value and confers an unfair productivity advantage; therefore, such works should not count toward T&P benchmarks. Other faculty members find AI valuable in their research and writing and advocate for permitting its scholarly use. Still others refuse to use or allow AI due to its societal and environmental impacts.

Scenario #7 – University Administrator Licensing of AI Services

You are a university administrator debating acquiring university-wide access to a popular AI model (e.g., *Copilot*, *ChatGPT*, *Gemini*, or *Claude*). You seek input on faculty guidance for its use and expense justification. Noting peer institutions promote their AI leadership through AI-related programs and grants, you wish to avoid falling behind.

Appendix C: Curated List of Resources

The AI and Research Working Group has curated a list of select resources that may serve as an entry point for those who are new to AI. The list is not exhaustive. Given that information about and development of AI is ongoing and constant, this list includes guidance from other scholarly societies and example resources where one might begin to dig deeper into each of the three areas of the Framework for AI Use in Language and Literary Scholarship.

MLA-CCCC Joint Task Force on Writing and AI Resources

- [*Working Paper 1: MLA-CCCC Joint Task Force on Writing and AI Working Paper: Overview of the Issues, Statement of Principles, and Recommendations*](#). July 2023.
- [*Working Paper 2: Generative AI and Policy Development: Guidance from the MLA-CCCC Task Force*](#). Apr. 2024.

- [*Working Paper 3: Building a Culture for Generative AI Literacy in College Language, Literature, and Writing*](#). Created by the MLA-CCCC Joint Task Force on Writing and AI and members of the Critical AI Literacy for Reading, Writing, and Languages Workshop. Oct. 2024.
- [“Initial Guidance for Evaluating the Use of AI in Scholarship and Creativity.”](#) 28 Jan. 2024.

Professional Organization Resources

- [*AAUP’s Artificial Intelligence and Academic Professions*](#)
- [*American Library Association’s Guidance on the Use of Artificial Intelligence in Libraries*](#) (draft)
- [APA Publishing’s “Principles for Responsible AI Use”](#)
- [*Viewfinder: A Toolkit for Values-Driven AI in Libraries and Archives*](#). Funded by the Institute of Museum and Library Services.

Additional Resources

- Lo, Leo S. [“AI Literacy: A Guide for Academic Libraries.”](#) *College and Research Libraries News*, vol. 86, no. 3, Mar. 2025.
- Murgia, Madhumita, et al. [“Generative AI Exists Because of the Transformer: This Is How It Works.”](#) *The Financial Times*, 12 Sept. 2023.
- Staiman, Avi. [“Why Authors Aren’t Disclosing AI Use and What Publishers Should \(Not\) Do About It.”](#) *The Scholarly Kitchen*, 27 Jan. 2026.
- Staiman, Avi. [“Why Authors Aren’t Disclosing AI Use and What Publishers Should \(Not\) Do About It, Part 2.”](#) *The Scholarly Kitchen*, 3 Feb. 2026.
- UNESCO. [Recommendation on the Ethics of Artificial Intelligence](#). 23 Nov. 2021.