

Last Words

Justify Your Prompts!

Eduardo Calò^{1*}, David M. Howcroft^{2*}, Leo Leppänen^{3*},
Saad Mahamood^{4*}, Simon Mille^{5*}, Patrícia Schmidtová^{6*},
and Emiel van Miltenburg^{7*,**}

¹Utrecht University

e.calo@uu.nl

²University of Aberdeen

david.howcroft@abdn.ac.uk

³University of Helsinki and University of Jyväskylä

leo.j.leppanen@jyu.fi

⁴Shopware

saad@saad.me.uk

⁵ADAPT Research Centre, Dublin City University

simon.mille@adaptcentre.ie

⁶Charles University

schmidtova@ufal.mff.cuni.cz

⁷Tilburg University

C.W.J.vanMiltenburg@tilburguniversity.edu

When you use a large language model (LLM) in your research, you often need to formulate a prompt to elicit some relevant output from the LLM. This step is challenging since (1) LLMs are known to be brittle and their results may vary drastically between different prompts; and (2) for any given task, there are infinitely many possible prompts. Thus we end up with the following problem: If you cannot try out infinitely many prompts, how do you justify your selected prompt? This article discusses different ways in which authors may motivate their choices, and urges authors to consider how their prompting strategy might be justified.

1. Introduction

The introduction of ChatGPT (OpenAI 2022) has clearly been a watershed moment. In the Before Times, whenever we wanted a language model to be able to perform a particular task, we would take a pretrained model and fine-tune it on a dataset exem-

* All authors contributed equally to this paper.

** Corresponding author.

Action Editor: Minlie Huang. Submission received: 5 February 2026; revised version received: ; accepted for publication: 14 March 2026.

<https://doi.org/10.1162/COLL.a.620>

© 2026 Association for Computational Linguistics

Published under a Creative Commons Attribution-NonCommercial-NoDerivatives 4.0 International (CC BY-NC-ND 4.0) license

plifying the task-to-be-performed. We knew that the generalizability of the model and our findings depended on the training data and the benchmark. Nowadays, however, a very large, pretrained, language model may be made to carry out the desired task through **prompting**: specifying the task to be carried out through natural language.^{1,2}

But what does this mean for the generalizability of our results? As the prompting search space is (effectively) infinite, any large language model (LLM) related research using prompts necessarily involves sampling: No study is able to evaluate the whole search space of prompts. This is not a new problem, however: Almost any research involves numerous *researcher degrees of freedom* which, if inadequately controlled for and documented, can lead researchers to false conclusions (Simmons, Nelson, and Simonsohn 2011; Wicherts et al. 2016). While much recent work focused on evaluating automatically generated texts has emphasized the importance of choosing the right automated metrics (Schmidtová et al. 2024) and human evaluation methods and providing adequate documentation to improve reproducibility (Belz et al. 2023), these are far from being the only degrees of freedom. Research using LLMs with prompting generally does not follow similar practices, despite the fact that an LLM's performance on a given task varies greatly with the chosen prompt. Subtle differences in prompts can dramatically impact LLM performance: changing prompt formatting (He et al. 2024), spelling (Schmidtová et al. 2025), phrasing (Lunardi et al. 2025), information structure. (Brucks and Toubia 2025), etc., all impact task performance. Beyond these issues, language itself is infinitely productive,³ with endless ways to ask a model to perform a particular task. This brittleness of performance and variability of valid prompts necessitate proper documentation of the prompt sampling and selection process in order for others to evaluate the reliability (and potential generalizability) of the study.

Discussions of how to decide and report on sampling strategies are plentiful across scientific fields.⁴ Perhaps most relevant for studies involving the creation of new software systems, the SIGSOFT reporting standard for software engineering research (Ralph et al. 2021) states that it is essential for any empirical study involving sampling to explain (a) the goal of the sampling, (b) the sampling strategy, (c) why the selected strategy is reasonable for the selected goal, and (d) the reasoning behind the selection of study objects. In this article, we (1) contextualize prompt justification as part of the larger metascientific discourse on reproducibility, (2) discuss different kinds of justifications that researchers might use and when these justifications may be appropriate, and (3) argue for the necessity of strong justifications and adequate documentation to ensure scientific rigor, improve reproducibility, and enable the collaborative development of generalizable findings. Our article contributes to a long history of metascientific discussion in the computational linguistics and natural language processing community. Earlier Last Words segments covered whether and how to compete with big tech (Kilgarrieff 2007), how to benefit and build upon research in other fields

1 The term “prompting” is overloaded; it has different meanings depending on the context. In this article, we consider “prompts” to be the initial instructions that tell the system what to do. For ease of presentation, we ignore the more complex case where systems consist of a model and a prompt, and users interact with that system through further prompts.

2 Although some research also considers embeddings as *continuous prompts*; see Liu et al. (2023) for discussion.

3 This relates as well to general problems with evaluating automatically generated texts: For all but the simplest cases there are many, many linguistically valid ways to complete a task (i.e., express the same meaning), which makes it impossible for a set of references to cover the full range of “correct” answers.

4 For example, in economics (Salle and Yıldızoğlu 2014), journalism (Firdaus, Aksar, and Gong 2024), medicine (Msaouel 2023), and psychology (Shen et al. 2011).

(Reiter 2007; Spärck Jones 2007; Wintner 2009; Krahmer 2010, particularly linguistics and psychology), reproducibility (Pedersen 2008), applications and impact (Belz 2009), the importance of negative results (Szpakowicz 2010), data ethics and integrity (Fort, Adda, and Cohen 2011), terminological clarity and epistemic soundness (Riezler 2014), and best practices for developing shared tasks (Nissim et al. 2017).⁵

2. Background

Before we consider prompt justification strategies, we first provide pointers to the literature on prompting strategies and prompting quality, after which we discuss related work in metascience for computational linguistics and the sampling justification strategies that are used in the psychology literature. These strategies form the inspiration for the prompt justification strategies listed herein.

2.1 Prompting Strategies

Many prompting strategies have been developed in the literature. Schulhoff et al. (2025) provide a comprehensive taxonomy, where 58 different techniques have been identified, organized in six main categories. A concurrent survey by Sahoo et al. (2025) found 41 different prompting techniques, organized by application domain. This gives prompt authors some sense of when to use which set of prompting techniques, although the number of options remains daunting. An earlier survey by Liu et al. (2023) offers a more modest taxonomy of prompting methods, but it is just constructed at a higher level of abstraction (e.g., by having “hand-crafted prompt templates” as the lowest level, and not distinguishing different ways to formulate a prompt). Finally, the survey presented in Vatsal, Dubey, and Singh (2025) focuses on multilingual prompt engineering, covering 39 prompting techniques applied to 30 multilingual NLP tasks.

Of course, these different techniques can also be combined in different ways, leading to a combinatorial explosion. So it is natural to ask: With all these prompting techniques in mind, how can we write a good prompt?

2.2 Prompt Quality

We are starting to see more and more evidence-based guidelines on prompt writing. For example, Long et al. (2025) carried out a thorough study on the effectiveness of different prompting techniques for different NLP tasks, resulting in a set of recommendations for prompt authors. Other guidelines include, for example, those provided by Maaz et al. (2024) and Geroimenko (2025). But even with recommendations or prompt assessments at hand, there are still *many* different ways to formulate your prompt. Thus the question arises: How do you justify the exact formulation of your prompt?

⁵ Other related studies focus on, for example, best practices for studying usability (Dybkjaer and Bernsen 2001) and developing language resources (Arranz et al. 2008) or how to adapt to changing methodologies (Saphra et al. 2024). Not to mention papers such as Church (2005), which do not address research practices but the practices that apply to publication and review.

3. Justification Strategies

As a general research principle, for any kind of study, researchers should always be able to motivate their design choices.⁶ We take our inspiration from the psychology literature, where studies have been published on ways to justify sample sizes (Lakens 2022) and alpha values for statistical significance testing (Lakens et al. 2018). Taking the former as an example, Lakens (2022) identifies different strategies that a researcher might use to justify their choice of sample size. We may paraphrase these as follows:

1. *Avoid the problem altogether* by using the entire population.
2. *Appeal to resource constraints* as a reason for a smaller sample.
3. *Compute an optimal solution* by using a power analysis or other statistical tools to obtain a sample size that is large enough to draw valid conclusions, but not larger than needed.⁷
4. *Use heuristics* such as sample sizes mentioned in the literature, or “rules of thumb” mentioned by experienced researchers.
5. *No justification*. The researcher honestly reports not to have any justification.

These justification strategies have clear parallels with the justification strategies that may be used to motivate prompting decisions. Applying their categorization scheme to this domain results in the justification types discussed below. (In practice, these different types may often be combined.)

3.1 Avoiding the Problem or Embracing the Variation

Avoiding the problem in the sense described by Lakens (2022)—by using the whole population—is not feasible. It is literally impossible to test every possible prompt, and any selection of prompts will by necessity involve sampling. However, researchers might still choose to “avoid” the problem by simply acknowledging it and using multiple prompts, documenting the chosen prompts and the resulting variation. This approach is advocated by Mizrahi et al. (2024), who compared the performance of 20 different LLMs across 39 tasks. Instead of using a single set of instructions, the authors used a large set of instruction paraphrases. They show that both manual and automatic paraphrases lead to a high degree of variation in model performance. Based on their results, the authors argue that LLMs should be evaluated using a large collection of prompts and reporting the *average performance*, documenting the range of potential

6 For example, the Netherlands Code of Conduct for Research Integrity (applicable to all researchers in the country) specifically instructs researchers to “Ensure that the methods you employ are well justified” (KNAW et al. 2018). More generally, The European Code of Conduct for Research Integrity requires researchers to “design, carry out, analyze, and document research in a careful, transparent, and well-considered manner” (ALLEA 2023, p. 7).

7 We have merged two items from the original study (*Accuracy* and *Power analysis*) since the distinction is not relevant to us here.

outcomes.^{8,9} Although this strategy is elegant in the way it embraces variation, it does run into three issues:

1. Because there are infinitely many possible prompts, you still need to make some kind of selection. Both the selection procedure (e.g., an initial prompt and the generation of paraphrases) and the sample size still need to be justified.
2. This approach is computationally expensive, since the number of computations linearly increases with the number of paraphrases.
3. Depending on the task, it may not be feasible to use multiple paraphrases, since the LLM outputs may be used for downstream processes, where the amount of effort (super)linearly increases with the number of outputs.

With these observations we do not wish to dispense with the multiple prompts approach; we only want to highlight the complexity of the situation. Critically analyzing this approach leads to interesting new questions, such as: How should one compose the set of prompts that is used for this approach,¹⁰ and how large should that set be?

3.2 Resource Constraints and Reasonable Effort

Resources like time and money are finite and as a natural consequence researchers must make understandable compromises.¹¹ Although it is hard to draw the line, authors might appeal to the *reasonable effort heuristic*: As a community, we expect the authors to put in a reasonable amount of effort to develop the prompts for their study (but it is not clear how much that should be, exactly). Generally speaking, the first prompt that comes to mind may not be optimal, but it is also not desirable to spend years developing the optimal prompt. So, it is important to show your work, describing the iterative process through which the prompts came about. If the results of the study hinge on the quality of the prompt, it is important to mention this as a limitation of the study.

3.3 Evidence-based Justification

To demonstrate whether a particular prompt leads to plausibly good results for a given LLM there needs to be some form of evidential proof. We have identified three particular subtypes of justifications that authors could use: showing that the results are good enough; showing that the prompts are optimal; and arguing for prompts on the basis of representativeness.

8 This also avoids researchers cherry-picking the results that favor their narrative.

9 As a side note, this approach also shows parallels with *multiverse analyzes* that are carried out in psychological research (e.g., Steegen et al. 2016), Mazei et al. 2025). The idea here is that, instead of carrying out only a single analysis, the researcher explore multiple different analyzes based on different reasonable assumptions (e.g., using different exclusion criteria, or different ways to discretize a continuous variable).

10 For example, could we use one LLM to generate a prompt for another LLM in the same model family (or even having LLMs generate prompts for themselves, e.g., as in Nye, Mee, and Core 2023).

11 For a discussion of barriers in NLP research, see van Miltenburg et al. (2023).

3.3.1 Results Are Good Enough. If the results using a particular prompt are good enough for your purposes, then it may not be relevant whether the prompts could have been improved. Of course, this shifts the problem from justification to demonstration: How can you demonstrate that the outputs are good enough? One option to consider is to carry out a *pilot study* where a representative number of outputs is assessed according to some relevant quality metric (Sripada, Reiter, and Hawizy 2005; van Miltenburg et al. 2021). This strategy may not be appropriate when the goal of your study is to show how good LLM outputs can be. But it may be all you need if the LLM outputs only form part of the research process, where you build on the outputs toward a larger goal.

3.3.2 Prompts Are (Locally) Optimal. If there is an objective way to assess the quality of the output, then it may be feasible to compute an optimal prompt for a given LLM-dataset combination. Indeed, the survey by Liu et al. (2023) discusses two general strategies that LLM researchers may use: automatically discovering discrete (Natural Language) or continuous (embeddings) prompts.¹² For example, Zhou et al. (2023) show for 21 (out of 24) different tasks that LLMs achieve comparable or better performance in formulating prompts than human annotators. Moreover, Do et al. (2025) present an approach to automatically select the optimal prompt from a finite set of synthetic candidate prompts, based on clustering and a prompt evaluator that scores the prompts and selects the best one. In theory, given the infinite search space, it may still be possible to formulate a better prompt, but this is the closest we can get to an optimal prompt. There are two clear limitations to this justification approach:

1. As noted above, this approach may not be feasible for tasks where the quality of the output is hard to assess.
2. This justification approach is only accessible to a select group of people with the technical abilities to set up the software to find the optimal prompt for the task at hand. Future researchers may want to find ways to reduce this barrier.

Because most optimization strategies offer *locally* optimal solutions, researchers should be careful in their justification not to imply that these solutions are also *globally* optimal (i.e., the best possible prompts).¹³

3.3.3 When Optimality is not Relevant: Selecting “Representative Prompts”. This strategy involves an appeal to the *ecological validity* of the study. When selecting prompts for a real-life setting, it is important to realize that *most people are not prompt experts*. So it is likely that the average user will provide suboptimal prompts (Zamfirescu-Pereira et al. 2023). If you are interested in the average-case performance of an LLM, you should also be using average prompts. These could be collected through a pilot study, where a representative set of potential users might interact with an LLM themselves, or where they are asked through a survey to write prompts that they might use for specified use cases.

¹² One potential downside of continuous prompts is that we lose some of the interpretability that comes with human-readable prompts.

¹³ At some point we just have to fall back on an appeal to resource constraints.

3.4 The “Prior Research” Heuristic

There are two ways in which prior research can be used to establish the prompt for a new study: using exactly the same prompt as in earlier research, or phrasing a prompt based on examples or suggestions from the literature.

3.4.1 Exactly the Same Prompt. You could copy the exact prompt from an earlier study that performed the same task. This *might* yield good results, but it remains important to consider the context in which the prompt is used:

- Prompt performance might differ between models; a prompt that is optimal for one model may perform worse than its competitors when used with a different model (as, e.g., demonstrated in Schmidová et al. 2025).
- Earlier results may have been carried out with a *different version* of the same model. It is not guaranteed that for a given prompt the performance will remain stable across different model iterations.
- Generalizability of the findings of earlier studies depends on the reliability of those studies. (If a prompt was used in prior research, that does not automatically mean it is a good prompt.)

3.4.2 Conceptually the Same Prompt. Instead of using exactly the same prompt as in earlier research, one might also use a prompt that is *conceptually similar*. By this we mean that the prompt is either:

1. Based on best practices reported in earlier studies (e.g., Long et al. 2025), or:
2. Based on successful examples from earlier studies solving similar challenges.

The same limitations as in §3.4.1 apply here. Moreover, because there are likely to be alternative prompt formulations, one might still need to appeal to resource constraints/reasonable effort (§3.2).

3.5 No Justification

Finally, as with Lakens’s (2022) list of justifications, one might honestly report not to have any justification for their choice of prompt. This might be seen as a weakness when you make strong claims about model performance, but different studies may have different goals. It does not always matter whether the authors are using “the perfect prompt” (if such a thing even exists).¹⁴ We will now turn to some example goals that influence prompt design.

¹⁴ And as a side note: Ideally we would work towards models that are robust to prompt variation anyway.

4. Research Goals Shaping Prompt Design

Roughly speaking, the role that LLMs play in a study determines what properties (and what importance) their prompts ought to have, and therefore what kind of justification is required. A key distinction is whether the study is exploratory or confirmatory in nature (Schwab and Held 2020). Exploratory studies investigate a space of prompting choices in order to discover phenomena or generate hypotheses, and therefore benefit from prompt variation and transparency about design iterations. Confirmatory studies aim to make precise claims about model behavior or performance, and therefore require a stronger justification of the prompt as a controlled experimental condition. With this in mind, we identify four broad perspectives under which LLMs are commonly investigated in NLP and adjacent fields.¹⁵

1. **LLMs as Tools to Help Users Achieve a Task.** The focus in this paradigm is on whether an LLM can assist or automate a practical goal such as, e.g., summarization (Chuang et al. 2024), tutoring (Nye, Mee, and Core 2023), planning (Sun et al. 2024), or domain-specific retrieval (Zhuang et al. 2024). Prompts should approximate realistic interaction: Most users are not prompt engineers, and their inputs may be brief or suboptimal. Studies should therefore justify prompts in terms of plausibility, user-centered design, or performance under natural variation, potentially including representative prompt sampling or prompts sourced from prior practice.
2. **LLMs as Objects of Study.** When probing model capabilities (Liang et al. 2022, Section 8.2) or behavior (Chang and Bergen 2024), prompts act as experimental stimuli. Small prompt changes can produce large differences in output, making it difficult to separate model properties from prompt artifacts. Here, exploratory prompt variation is often necessary to discover how robust a finding is. Stronger prompt justification is required when making confirmatory claims about what a model can or cannot reliably do.
3. **LLMs as Measurement Instruments.** LLMs are increasingly used to rate or classify content (Li et al. 2024), as, e.g., judges of quality (Sedoc et al. 2025), factuality (Zhang and Gao 2023), or safety (Yuan et al. 2025). In such cases, prompts define the measurement protocol: what is being judged, according to which criteria, and in what format. Justification must therefore address validity (alignment with the intended construct), reliability (consistency across rephrasings), and transparency (reproducibility of the evaluation conditions).
4. **LLMs as Sociotechnical Actors.** In deployed applications—educational platforms (Chu et al. 2025), public services (Musumeci et al. 2024), creative tools (Dang et al. 2022), or decision-support pipelines (Eigner and Händler 2024)—prompt choices shape not only output quality but also users and institutions. Under this perspective, justification extends

¹⁵ These perspectives are not exhaustive nor mutually exclusive, but they capture the dominant motivations behind prompting choices in current research.

to social consequences: safety, bias, robustness, and governance. Prompts generally encode normative assumptions about appropriate behavior,¹⁶ and such assumptions should be made explicit.

Prompt choice is, therefore, both goal-dependent and tied to the evidentiary standard a study intends to meet, which in turn determines how rigorously prompts must be justified.

5. Future Research

It is too early to talk about “best practices” in prompt justification without the existence of an empirical evidence base. Although we believe that prompt development needs reporting standards, it is still unclear what those should look like, exactly. In the near term, it is useful for researchers to document their practices to enable future meta-scientific literature surveys to codify best practices. To this end, we offer some initial recommendations. If you do not do so already, a good first step is to keep a lab notebook to record your prompting history and the motivation behind each change. Best practices for keeping lab notebooks are discussed, for example, by Pain (2019) and Higgins, Nogiwa-Valdez, and Stevens (2022). When reporting your study, you should be able to answer the following questions:

- What kind of **justification** philosophy are you using and why?
- What **techniques** did you use to develop your prompt(s)?
- What **set of candidate prompts** did you select from?
- What **criteria** did you use to select your prompt(s)?
- How did candidate prompts **perform** with respect to those criteria?
- What **prompt(s) did you end up using?**

Capturing this information provides a minimum starting point, enabling readers to appreciate the generalizability of your findings and to understand how they relate to their own and others’ research.

6. Final Note

At least for the foreseeable future, prompts are here to stay. The “naturalness” of this natural language interface is a double-edged sword: On the one hand, it made state-of-the-art technology accessible to the masses, but on the other hand, its brittleness poses serious risks to the generalizability of our results. Not just because small variations in the input can lead to serious performance degradation, but also because the intuitive interface makes it easy to forget the vast space of input parameters that we are manipulating through natural language. Thus, we need to make a conscious effort to make our prompting strategies explicit. As we have seen, earlier work has issued a call to provide justifications for all of our design choices; let’s start by justifying our prompts!

¹⁶ These assumptions might either be explicit (stating how the LLM should behave) or implicit (omitting such instructions and relying on the default behavior of the model).

Acknowledgments

D.M.H. was supported by CRUK grant EDDPJT-May23/100001. S.Mi. was funded by the Irish Department of Tourism, Culture, Arts, Gaeltacht, Sport and Media via the eSTÓR project, and he benefits from being a member of the SFI Ireland funded ADAPT Research Centre. P.S. was supported by Charles University projects GAUK 252986 and SVV project 260 821.

References

- ALLEA. 2023. *The European Code of Conduct for Research Integrity – Revised edition 2023*. ALLEA — All European Academies. <https://doi.org/10.26356/ECOC>
- Arranz, Victoria, Franck Gandcher, Valérie Mapelli, and Khalid Choukri. 2008. A guide for the production of reusable language resources. In *Proceedings of the Sixth International Conference on Language Resources and Evaluation (LREC'08)*. <https://doi.org/10.63317/2xe4npnzx3cs>
- Belz, Anja. 2009. Last words: That's nice ... what can you do with it? *Computational Linguistics*, 35(1):111–118. <https://doi.org/10.1162/coli.2009.35.1.111>
- Belz, Anya, Craig Thomson, Ehud Reiter, Gavin Abercrombie, Jose M. Alonso-Moral, Mohammad Arvan, Anouck Braggaar, Mark Cieliebak, Elizabeth Clark, Kees van Deemter, et al. 2023. Missing information, unresponsive authors, experimental flaws: The impossibility of assessing the reproducibility of previous human evaluations in NLP. In *Proceedings of the Fourth Workshop on Insights from Negative Results in NLP*, pages 1–10. <https://doi.org/10.18653/v1/2023.insights-1.1>
- Brucks, Melanie and Olivier Toubia. 2025. Prompt architecture induces methodological artifacts in large language models. *PLOS ONE*, 20(4):1–13. <https://doi.org/10.1371/journal.pone.0319159>, PubMed: 40293988
- Chang, Tyler A. and Benjamin K. Bergen. 2024. Language model behavior: A comprehensive survey. *Computational Linguistics*, 50(1):293–350. https://doi.org/10.1162/coli_a_00492
- Chu, Zhendong, Shen Wang, Jian Xie, Tinghui Zhu, Yibo Yan, Jinheng Ye, Aoxiao Zhong, Xuming Hu, Jing Liang, Philip S. Yu, et al. 2025. LLM agents for education: Advances and applications. *arXiv preprint arXiv:2503.11733*. <https://doi.org/10.18653/v1/2025.findings-emnlp.743>
- Chuang, Yu Neng, Ruixiang Tang, Xiaolian Jiang, and Xia Hu. 2024. SPeC: A soft prompt-based calibration on performance variability of large language model in clinical notes summarization. *Journal of Biomedical Informatics*, 151:104606. <https://doi.org/10.1016/j.jbi.2024.104606>, PubMed: 38325698
- Church, Kenneth. 2005. Last words: Reviewing the reviewers. *Computational Linguistics*, 31(4):575–578. <https://doi.org/10.1162/089120105775299131>
- Dang, Hai, Lukas Mecke, Florian Lehmann, Sven Goller, and Daniel Buschek. 2022. How to prompt? Opportunities and challenges of zero-and few-shot learning for human-AI interaction in creative applications of generative models. In *Generative AI and HCI Workshop at CHI '22*.
- Do, Viet Tung, Xuan-Quang Nguyen, Van-Khanh Hoang, Duy-Hung Nguyen, Shahab Sabahi, Jeff Yang, Hajime Hotta, Minh-Tien Nguyen, and Hung Le. 2025. Automatic prompt selection for large language models. In *Advances in Knowledge Discovery and Data Mining*, pages 91–102. https://doi.org/10.1007/978-981-96-8180-8_8
- Dybkjær, Laila and Niels O. Bernsen. 2001. Usability evaluation in spoken language dialogue systems. In *Proceedings of the ACL 2001 Workshop on Evaluation Methodologies for Language and Dialogue Systems*. <https://doi.org/10.3115/1118053.1118055>
- Eigner, Eva and Thorsten Händler. 2024. Determinants of LLM-assisted decision-making. *arXiv preprint arXiv:2402.17385*.
- Firdaus, Amira, Iffat Ali Aksar, and Jiankun Gong. 2024. Small world sampling: Qualitative sample reliability and validity for efficient and effective recruitment of journalists as research participants. *Journalism*, 25(2):466–485. <https://doi.org/10.1177/14648849221124997>
- Fort, Karén, Gilles Adda, and K. Bretonnel Cohen. 2011. Last words: Amazon Mechanical Turk: Gold mine or coal mine? *Computational Linguistics*, 37(2):413–420. https://doi.org/10.1162/coli_a_00057
- Geroimenko, Vladimir. 2025. Key principles of good prompt design. In *The Essential Guide to Prompt Engineering: Key Principles, Techniques, Challenges, and Security Risks*, pages 17–36. Springer Nature Switzerland.

- https://doi.org/10.1007/978-3-031-86206-9_2
- He, Jia, Mukund Rungta, David Koleczek, Arshdeep Sekhon, Franklin X. Wang, and Sadid Hasan. 2024. Does prompt formatting have any impact on LLM performance? *arXiv preprint arXiv:2411.10541*.
- Higgins, Stuart G., Akemi A. Nogiwa-Valdez, and Molly M. Stevens. 2022. Considerations for implementing electronic laboratory notebooks in an academic research environment. *Nature Protocols*, 17(2):179–189. <https://doi.org/10.1038/s41596-021-00645-8>, PubMed: 35031789
- Kilgariff, Adam. 2007. Last words: Googleology is bad science. *Computational Linguistics*, 33(1):147–151. <https://doi.org/10.1162/coli.2007.33.1.147>
- KNAW, NFW, NWO, TO2-federatie, Vereniging Hogescholen, and VSNU. 2018. Nederlandse gedragscode wetenschappelijke integriteit. <https://doi.org/10.17026/DANS-2CJ-NVWU>
- Krahmer, Emiel. 2010. Last words: What computational linguists can learn from psychologists (and vice versa). *Computational Linguistics*, 36(2):285–294. <https://doi.org/10.1162/coli.2010.36.2.36201>
- Lakens, Daniël. 2022. Sample size justification. *Collabra: Psychology*, 8(1):33267. <https://doi.org/10.1525/collabra.33267>
- Lakens, Daniel, Federico G. Adolphi, Casper J. Albers, Farid Anvari, Matthew A. J. Apps, Shlomo E. Argamon, Thom Baguley, Raymond B. Becker, Stephen D. Benning, Daniel E. Bradford, et al. 2018. Justify your alpha. *Nature Human Behaviour*, 2(3):168–171. <https://doi.org/10.1038/s41562-018-0311-x>
- Li, Haitao, Qian Dong, Junjie Chen, Huixue Su, Yujia Zhou, Qingyao Ai, Ziyi Ye, and Yiqun Liu. 2024. LLMs-as-judges: A comprehensive survey on LLM-based evaluation methods. *arXiv preprint arXiv:2412.05579*.
- Liang, Percy, Rishi Bommasani, Tony Lee, Dimitris Tsipras, Dilara Soylu, Michihiro Yasunaga, Yian Zhang, Deepak Narayanan, Yuhuai Wu, Ananya Kumar, et al. 2022. Holistic evaluation of language models. *arXiv preprint arXiv:2211.09110*.
- Liu, Pengfei, Weizhe Yuan, Jinlan Fu, Zhengbao Jiang, Hiroaki Hayashi, and Graham Neubig. 2023. Pre-train, prompt, and predict: A systematic survey of prompting methods in natural language processing. *ACM Computing Surveys*, 55(9). <https://doi.org/10.1145/3560815>
- Long, Do Xuan, Duy Dinh, Ngoc-Hai Nguyen, Kenji Kawaguchi, Nancy F. Chen, Shafiq Joty, and Min-Yen Kan. 2025. What makes a good natural language prompt? In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 5835–5873. <https://doi.org/10.18653/v1/2025.acl-long.292>
- Lunardi, Riccardo, Vincenzo Della Mea, Stefano Mizzaro, and Kevin Roitero. 2025. On robustness and reliability of benchmark-based evaluation of LLMs. In *Proceedings of the 28th European Conference on Artificial Intelligence (ECAI)*, Frontiers in Artificial Intelligence and Applications, volume 413, pages 4603–4610. <https://doi.org/10.3233/FAIA251363>
- Maaz, Sara, Janice C. Palaganas, Gerry Palaganas, and Maria Bajwa. 2024. A guide to prompt design: Foundations and applications for healthcare simulationists. *Frontiers in Medicine (Lausanne)*, 11:1504532. <https://doi.org/10.3389/fmed.2024.1504532>, PubMed: 39980724
- Mazei, Jens, Cort W. Rudolph, Hannes Zacher, and Joachim Hüffmeier. 2025. Do not put all of your eggs in one basket: Multiverse analysis in applied psychology. *Journal of Applied Psychology*, 110(11):1511–1537. <https://doi.org/10.1037/ap10001291>, PubMed: 40323882
- Mizrahi, Moran, Guy Kaplan, Dan Malkin, Rotem Dror, Dafna Shahaf, and Gabriel Stanovsky. 2024. State of what art? A call for multi-prompt LLM evaluation. *Transactions of the Association for Computational Linguistics*, 12:933–949. https://doi.org/10.1162/tac1_a-00681
- Msaouel, Pavlos. 2023. The role of sampling in medicine. *Harvard Data Science Review*, 5(3). <https://hdsr.mitpress.mit.edu/pub/s9cynk9u>, <https://doi.org/10.1162/99608f92.bc6818d3>
- Musumeci, Emanuele, Michele Brienza, Vincenzo Suriani, Daniele Nardi, and Domenico Daniele Bloisi. 2024. LLM based multi-agent generation of semi-structured documents from semantic templates in the public administration domain. In *International Conference on Human-Computer Interaction*, pages 98–117. https://doi.org/10.1007/978-3-031-60615-1_7
- Nissim, Malvina, Lasha Abzianidze, Kilian Evang, Rob van der Goot, Hessel Haagsma, Barbara Plank, and Martijn

- Wieling, 2017. Last words: Sharing is caring: The future of shared tasks. *Computational Linguistics*, 43(4):897–904. https://doi.org/10.1162/COLI_a.00304
- Nye, Benjamin D., Dillon Mee, and Mark G. Core. 2023. Generative large language models for dialog-based tutoring: An early consideration of opportunities and concerns. In *LLM@ AIED*, pages 78–88.
- OpenAI. 2022. ChatGPT – Introducing ChatGPT. Company blogpost: <https://openai.com/index/chatgpt/>. Accessed 19 December 2025.
- Pain, Elizabeth. 2019. How to keep a lab notebook. <https://www.science.org/content/article/how-keep-lab-notebook>. Accessed 19 March 2026.
- Pedersen, Ted. 2008. Last words: Empiricism is not a matter of faith. *Computational Linguistics*, 34(3):465–470. <https://doi.org/10.1162/coli.2008.34.3.465>
- Ralph, Paul, Nauman bin Ali, Sebastian Baltes, Domenico Bianculli, Jessica Diaz, Yvonne Dittrich, Neil Ernst, Michael Felderer, Robert Feldt, Antonio Filieri, et al. 2021. Empirical standards for software engineering research. *arXiv preprint arXiv:2010.03525*.
- Reiter, Ehud. 2007. Last words: The shrinking horizons of computational linguistics. *Computational Linguistics*, 33(2):283–287. <https://doi.org/10.1162/coli.2007.33.2.283>
- Riezler, Stefan. 2014. Last words: On the problem of theoretical terms in empirical computational linguistics. *Computational Linguistics*, 40(1):235–245. https://doi.org/10.1162/COLI_a.00182
- Sahoo, Pranab, Ayush Kumar Singh, Sriparna Saha, Vinija Jain, Samrat Mondal, and Aman Chadha. 2025. A systematic survey of prompt engineering in large language models: Techniques and applications. *arXiv preprint arXiv:2402.07927*.
- Salle, Isabelle and Murat Yıldızoglu. 2014. Efficient sampling and meta-modeling for computational economic models. *Computational Economics*, 44(4):507–536. <https://doi.org/10.1007/s10614-013-9406-7>
- Saphra, Naomi, Eve Fleisig, Kyunghyun Cho, and Adam Lopez. 2024. First tragedy, then parse: History repeats itself in the new era of large language models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (Volume 1: Long Papers)*, pages 2310–2326. <https://doi.org/10.18653/v1/2024.naacl-long.128>
- Schmidtová, Patrícia, Niyati Bafna, Seth Aycok, Gianluca Vico, Wiktor Kamzela, Katharina Hämmerl, and Vilém Zouhar. 2025. How important is ‘perfect’ English for machine translation prompts? *arXiv preprint arXiv:2507.09509*. <https://doi.org/10.18653/v1/2026.findings-eacl.38>
- Schmidtová, Patrícia, Saad Mahamood, Simone Balloccu, Ondrej Dusek, Albert Gatt, Dimitra Gkatzia, David M. Howcroft, Ondrej Platek, and Adarsa Sivaprasad. 2024. Automatic metrics in natural language generation: A survey of current evaluation practices. In *Proceedings of the 17th International Natural Language Generation Conference*, pages 557–583. <https://doi.org/10.18653/v1/2024.inlg-main.44>
- Schulhoff, Sander, Michael Ilie, Nishant Balepur, Konstantine Kahadze, Amanda Liu, Chenglei Si, Yinheng Li, Aayush Gupta, Hyojung Han, Sevien Schulhoff, et al. 2025. The prompt report: A systematic survey of prompt engineering techniques. Available <https://trigaten.github.io/PromptSurvey.Site/>.
- Schwab, Sophie and Leonhard Held. 2020. Different worlds: Confirmatory versus exploratory research. *Significance*, 17(2):8–9. <https://doi.org/10.1111/1740-9713.01369>
- Sedoc, João, Simon Mille, Miruna Adriana Clinciu, Yixin Liu, Kaustubh Dhole, and Saad Mahamood. 2025. The 2024 GEM shared task on multilingual data-to-text generation: English and Spanish qualitative evaluation results. In *Proceedings of the 18th International Natural Language Generation Conference: Generation Challenges*, pages 1–36.
- Shen, Winny, Thomas B. Kiger, Stacy E. Davies, Rena L. Rasch, Kara M. Simon, and Deniz S. Ones. 2011. Samples in applied psychology: Over a decade of research in review. *Journal of Applied Psychology*, 96(5):1055–1064. <https://doi.org/10.1037/a0023322>, PubMed: 21463013
- Simmons, Joseph P., Leif D. Nelson, and Uri Simonsohn. 2011. False-positive psychology: Undisclosed flexibility in data collection and analysis allows presenting

- anything as significant. *Psychological Science*, 22(11):1359–1366. <https://doi.org/10.1177/0956797611417632>, PubMed: 22006061
- Spärck Jones, Karen. 2007. Last words: Computational linguistics: What about the linguistics? *Computational Linguistics*, 33(3):437–441. <https://doi.org/10.1162/coli.2007.33.3.437>
- Sripada, Somayajulu, Ehud Reiter, and Lezan Hawizy. 2005. Evaluation of an NLG system using post-edit data: Lessons learnt. In *Proceedings of the Tenth European Workshop on Natural Language Generation (ENLG-05)*.
- Steegen, Sara, Francis Tuerlinckx, Andrew Gelman, and Wolf Vanpaemel. 2016. Increasing transparency through a multiverse analysis. *Perspectives on Psychological Science*, 11(5):702–712. <https://doi.org/10.1177/1745691616658637>, PubMed: 27694465
- Sun, Jingkai, Qiang Zhang, Yiqun Duan, Xiaoyang Jiang, Chong Cheng, and Renjing Xu. 2024. Prompt, plan, perform: LLM-based humanoid control via quantized imitation learning. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, pages 16236–16242. <https://doi.org/10.1109/ICRA57147.2024.10610948>
- Szpakowicz, Stan. 2010. Last words: Failure is an orphan (let’s adopt). *Computational Linguistics*, 36(1):157–158. <https://doi.org/10.1162/coli.2010.36.1.36105>
- van Miltenburg, Emiel, Miruna Clinciu, Ondřej Dušek, Dimitra Gkatzia, Stephanie Inglis, Leo Leppänen, Saad Mahamood, Emma Manning, Stephanie Schoch, Craig Thomson, and Luou Wen. 2021. Underreporting of errors in NLG output, and what to do about it. In *Proceedings of the 14th International Conference on Natural Language Generation*, pages 140–153. <https://doi.org/10.18653/v1/2021.inlg-1.14>
- van Miltenburg, Emiel, Miruna Clinciu, Ondřej Dušek, Dimitra Gkatzia, Stephanie Inglis, Leo Leppänen, Saad Mahamood, Stephanie Schoch, Craig Thomson, and Luou Wen. 2023. Barriers and enabling factors for error analysis in NLG research. *Northern European Journal of Language Technology*, 9. <https://doi.org/10.3384/nejlt.2000-1533.2023.4529>
- Vatsal, Shubham, Harsh Dubey, and Aditi Singh. 2025. Multilingual prompt engineering in large language models: A survey across NLP tasks. ArXiv:2505.11665 [cs]. <https://doi.org/10.48550/arXiv.2505.11665>, <https://doi.org/10.1109/ACCESS.2026.3702852>
- Wicherts, Jelte M., Coosje L. S. Veldkamp, Hilde E. M. Augusteijn, Marjan Bakker, Robbie C. M. Van Aert, and Marcel A. L. M. Van Assen. 2016. Degrees of freedom in planning, running, analyzing, and reporting psychological studies: A checklist to avoid p-hacking. *Frontiers in Psychology*, 7. <https://doi.org/10.3389/fpsyg.2016.01832>, PubMed: 27933012
- Wintner, Shuly. 2009. Last words: What science underlies natural language engineering? *Computational Linguistics*, 35(4):641–644. <https://doi.org/10.1162/coli.2009.35.4.35409>
- Yuan, Xiaohan, Jinfeng Li, Dongxia Wang, Yuefeng Chen, Xiaofeng Mao, Longtao Huang, Jialuo Chen, Hui Xue, Xiaoxia Liu, Wenhai Wang, et al. 2025. S-eval: Towards automated and comprehensive safety evaluation for large language models. *Proceedings of the ACM on Software Engineering*, 2(ISSTA):2136–2157. <https://doi.org/10.1145/3728971>
- Zamfirescu-Pereira, J. D., Richmond Y. Wong, Bjoern Hartmann, and Qian Yang. 2023. Why Johnny can’t prompt: How non-AI experts try (and fail) to design LLM prompts. In *Proceedings of the 2023 CHI Conference on Human Factors in Computing Systems*, CHI ’23. <https://doi.org/10.1145/3544548.3581388>
- Zhang, Xuan and Wei Gao. 2023. Towards LLM-based fact verification on news claims with a hierarchical step-by-step prompting method. *arXiv preprint arXiv:2310.00305*. <https://doi.org/10.18653/v1/2023.ijcnlp-main.64>
- Zhou, Yongchao, Andrei Ioan Muresanu, Ziwen Han, Keiran Paster, Silviu Pitis, Harris Chan, and Jimmy Ba. 2023. Large language models are human-level prompt engineers. In *The Eleventh International Conference on Learning Representations*.
- Zhuang, Shengyao, Xueguang Ma, Bevan Koopman, Jimmy Lin, and Guido Zuccon. 2024. Promptreps: Prompting large language models to generate dense and sparse representations for zero-shot document retrieval. *arXiv preprint arXiv:2404.18424*. <https://doi.org/10.18653/v1/2024.emnlp-main.250>